ARYAN SIRDESAI

Generative AI Engineer | Agentic AI | LLM Systems | RAG Pipelines
+91 9673095364 | aryansirdesai@gmail.com | LinkedIn | GitHub

Generative AI Engineer with 2 years of end-to-end production experience - from embedding pipelines to deployed multi-tenant platforms. Specialise in Agentic AI systems, LangGraph-based orchestration, and RAG architectures that go beyond prototypes into real business infrastructure. Comfortable owning the full delivery: model integration, FastAPI backend, GCP deployment, and CI/CD.

SKILLS

Generative AI: LangChain, LlamaIndex, LangGraph, crewAI, AutoGen, RAG, Agentic RAG, MCP, A2A, Prompt Engineering (Types), LangSmith, RAGAS
LLMs: GPT-4, Gemini, Llama, Mistral, HuggingFace Transformers, Fine-Tuning (PEFT/LoRA)
Machine Learning: XGBoost, scikit-learn, PyTorch, SHAP, Feature Engineering, Predictive Modelling
Vector DBs: Qdrant, FAISS, Pinecone, ChromaDB, Milvus, MongoDB Atlas, PgVector, ElasticSearch
Backend / Cloud: FastAPI (ASGI), Flask (WSGI), Docker, JWT (RS256), GCP (Vertex AI, Cloud Run), AWS Bedrock, SageMaker
Databases & Tools: PostgreSQL, MongoDB, MySQL, SQLite | HuggingFace, Groq, OpenRouter, Celery, Redis, Kafka, Postman
Languages: Python 3, SQL

WORK EXPERIENCE

AI/ML Engineer (SDE-1) • Onelab Ventures Pvt. Ltd. | Pune — Mar 2025 – Apr 2026
Stack: *Python, FastAPI, XGBoost, LangGraph, crewAI, Vertex AI, GCP, AWS, Docker, JWT*

- Developed churn and product-interest prediction models using XGBoost with SHAP-based explainability; served via FastAPI microservices processing 600K+ user records.
- Designed white-label, multi-tenant platform architecture on GCP (Cloud Run + Vertex AI) enabling configurable deployments for multiple financial clients.
- Implemented JWT RS256 authentication and end-to-end CI/CD pipelines with Docker for containerised production deployments.

Data Science Intern • Shyena Tech Yarns Pvt. Ltd. | Pune — Jul 2024 – Mar 2025
Stack: *Python, FastAPI, LangChain, LlamaIndex, OpenAI API, Vector DBs, GCP*

- Built and deployed LLM-powered RAG applications for domain-specific Q&A, improving retrieval accuracy by 25–40% via optimised chunking strategies and embedding tuning.
- Engineered vector retrieval pipelines (FAISS, Pinecone) with prompt engineering to reduce hallucinations; exposed via FastAPI into client-facing production workflows.

PROJECTS

FinanSee – AI Financial Advisory Platform Live ↗

- End-to-end AI platform: behavioral bias detection, XGBoost-based user segmentation, and Agentic RAG for personalised financial recommendations deployed on GCP, multi-tenant, white-label.

AI Analyst – Natural Language to SQL GitHub ↗

- RAG Text-to-SQL system (LlamaIndex + GPT-3.5 + Chain-of-Thought) with auto SQL validation; Flask API enabling natural language queries over production databases.

Highly Customisable AI Assistant GitHub ↗

- Org-agnostic conversational AI adapting to any company's data in real-time — multi-format ingestion (PDF, JSON) with semantic search via RAG.

MCP Sentiment Analysis Server GitHub ↗

- Production-ready MCP server integrated with Claude Desktop via stdio transport — real-time sentiment analysis (polarity, subjectivity) using TextBlob; built with FastMCP + async/await architecture.

CERTIFICATIONS

Agents Course: Fundamentals of Agents — Hugging Face | AI Agents, ReAct, Tool Use, Memory, Multi-Agent Systems, smolagents — 2026
MCP Course: Fundamentals — Hugging Face | MCP Architecture, Servers, Tools, Claude Desktop Integration, A2A — 2025
LLM Course — Hugging Face | Transformers, Fine-tuning (PEFT / LoRA), RLHF, RAG, LLM Agents — 2024

EDUCATION

B.E. Computer Engineering — Dr. D. Y. Patil College of Engineering, Pune — 2020 – 2024
CGPA: 8.54